

SHOUREN WANG

☎ +1 (929) 302-1790 ✉ sxw992@case.edu 🌐 [Personal Website](#)
📄 [LinkedIn](#) 📄 [GitHub](#) 🎓 [Google Scholar](#)

EDUCATION

Case Western Reserve University

Ph.D. in Computer Science

Cleveland, OH

Aug. 2024 – Present

New York University

M.S. in Computer Engineering | GPA: 3.93/4.0

Brooklyn, NY

Sep. 2021 – May. 2023

Hunan University

Undergraduate in Software Engineering

*Undergraduate and B.E. in Digital Media Technology | GPA: 83.16/100.0
(A specialization of Software Engineering focusing on visual software/algorithm development)*

Changsha, China

Sep. 2015 – June 2017

June. 2017 – June 2019

EXPERIENCE

Research Intern

NYU Game Innovation Lab

Aug. 2023 – July 2024

Brooklyn, NY

- Explored the research and methods for creative AI for video games.
- Developed enjoyment focused game-playing agents for Street Fighter 2 via an extended PPO model and self-play.
- Advisor: [Julian Togelius](#), Professor, New York University.

Research and Development Engineer

AsiaInfo Technology

Sep 2019 – Sep 2020

Nanjing, China

- Tested and enhanced the performance of CTDI project as a QA group member, contributed to project efficiency.
- Manager: Ted Wang.

Teaching Assistant

Case Western Reserve University

Sep. 2024 – May 2026

Cleveland, OH

- Teaching assistant for *Large Language Models* and *Database Systems*: graded assignments and held office hours.
- Professors: [Xiaotian Han](#) (Large Language Models); [Yinghui Wu](#) (Database Systems).

PROJECTS

Agent Team Work Zone [\[GitHub\]](#) | *LLM Agents, Claude Code, Multi-Agent Workflow*

Apr. 2026 – Present

- Built and maintain an open-source persistence layer for Claude Code's Agent Teams, whose teammates vanish on session restart with no native recovery.
- Designed per-agent workstations (role, working context, checkpoints, commitments) and a team registry, so one command rebuilds an interrupted team from disk.
- Shipped 13 slash-command skills and hook-driven auto-checkpointing across four public releases, maintaining parallel editions through a breaking Claude Code API change.
- Dogfood the layer end-to-end: the agent team that develops it runs on it.

Path-Lock Expert: Efficient Hybrid Thinking Architecture [\[Paper\]](#) | *LLMs, Reasoning* Jan. 2026 – Jul. 2026

- *Accepted at COLM 2026.*
- Developed a parameter-efficient architecture approach to improve hybrid-thinking *mode controllability* (thinking vs. fast) without compromising overall task performance. Built an efficient training pipeline with curated data.
- Experiments on reasoning benchmarks (AIME24, MATH500, GPQA-Diamond) indicate improved mode separation/controllability with competitive accuracy, and greatly improved accuracy for no-think mode.

Demystifying Hybrid Thinking [\[Paper\]](#) | *LLMs, Reasoning, Post-Training*

Sep. 2025 – Feb 2026

- *Accepted at EMNLP 2026 Main.*
- Systematically analyzed the controllability of hybrid-thinking LLMs (think vs. no-think), showing incomplete mode separation and quantifying “reasoning leakage” via output length and reflection-token counts across multiple evaluations (MATH500/AIME24/GPQA/MMLU-STEM).
- Conducted controlled SFT ablations on data scale (20k–140k), pairing strategy, think:no-think ratio, and two-phase training; proposed a practical training recipe that preserves accuracy while cutting no-think verbosity (1085→585 tokens on MATH500) and reflection-token frequency (5917→522).

- *Accepted at ACL 2026 Findings.*
- Identified via attention analysis that hybrid-thinking mode switching is dominated by a few *token-level triggers*—“Okay” activates reasoning while `</think>\n\n` suppresses it—rather than high-level Think/No-think instructions.
- Proposed **Mid-Think**, a training-free prompting format combining both triggers to achieve Pareto-optimal accuracy–length trade-offs (MATH500: 92.1% acc. at 2589 tokens vs. 94.4% at 4904), outperforming fixed-token and prompt-based baselines across Qwen3-4B/8B/14B/32B.
- Applied Mid-Think to GRPO-based RL after SFT, cutting training time $\sim 15\%$ (54h $\rightarrow$ 46h) while improving Qwen3-8B on AIME (69.8% $\rightarrow$ 72.4%) and GPQA (58.5% $\rightarrow$ 61.1%).

PUBLICATIONS

*Equal contribution. †Corresponding author(s).

[1] **Path-Lock Expert: Separating Reasoning Mode in Hybrid Thinking via Architecture-Level Separation**

COLM 2026

- Authors: **Shouren Wang***, Wang Yang*, Chuang Ma, Debargha Ganguly, Vikash Singh, Chaoda Song, Xinpeng Li, Xianxuan Long, Vipin Chaudhary†, Xiaotian Han†
- Links: [\[Paper\]](#), [\[Repository\]](#)

[2] **Demystifying Hybrid Thinking: Can LLMs Truly Switch Between Think and No-Think?**

EMNLP 2026 Main

- Authors: **Shouren Wang***, Wang Yang*, Xianxuan Long, Qifan Wang, Vipin Chaudhary†, Xiaotian Han†
- Links: [\[Paper\]](#), [\[Repository\]](#)

[3] **Mid-Think: Training-Free Intermediate-Budget Reasoning via Token-Level Triggers**

ACL 2026 Findings

- Authors: Wang Yang*, **Shouren Wang***, Debargha Ganguly, Xinpeng Li, Chaoda Song, Vikash Singh, Vipin Chaudhary†, Xiaotian Han†
- Links: [\[Paper\]](#), [\[Repository\]](#)

[4] **A Survey on Efficient Protein Language Models**

TMLR 2026

- Authors: **Shouren Wang**, Debargha Ganguly, Vinooth Rao Kulkarni, Wang Yang, Zhuoran Qiao, Daniel Blankenberg, Vipin Chaudhary†, Xiaotian Han†
- Links: [\[Paper\]](#), [\[Repository\]](#)

[5] **When Domains Interact: Asymmetric and Order-Sensitive Cross-Domain Effects in Reinforcement Learning for Reasoning**

Under Review

- Authors: Wang Yang, **Shouren Wang**, Chaoda Song, Chuang Ma, Xinpeng Li, Nengbo Wang, Kaixiong Zhou, Vipin Chaudhary†, Xiaotian Han†
- Links: [\[Paper\]](#), [\[Repository\]](#)

[6] **A Survey on Agent Skills for LLMs: A Lifecycle Perspective from Construction to Ecosystems**

Under Review

- Authors: Wang Yang, Chaoda Song, Xinpeng Li, **Shouren Wang**, Nengbo Wang, Yanyan Zhang, Chuang Ma, Debargha Ganguly, Vikash Singh, Shuai Xu, Jing Ma, Yu Yin, Vipin Chaudhary†, Xiaotian Han†
- Links: [\[Paper\]](#), [\[Repository\]](#)

[7] **AgentCE-Bench: Agent Configurable Evaluation with Scalable Horizons and Controllable Difficulty under Lightweight Environments**

Under Review

- Authors: Wang Yang, Chaoda Song, Xinpeng Li, Debargha Ganguly, Chuang Ma, **Shouren Wang**, Zhihao Dou, Yuli Zhou, Vipin Chaudhary†, Xiaotian Han†
- Links: [\[Paper\]](#), [\[Repository\]](#)

[8] **MemTrace: Probing What Final Accuracy Misses in Long-Term Memory**

Under Review

- Authors: Xianxuan Long, Zhikai Chen, Shenglai Zeng, **Shouren Wang**, Kai Guo, Jiliang Tang[†]
- Links: [\[Paper\]](#)

[9] **Reliability-Gated Source Anchoring for Continual Test-Time Adaptation**

Under Review

- Authors: Vikash Singh, Debargha Ganguly, Weicon Chen, Sabyasachi Sahoo, Sreehari Sankar, Biyou Zhang, Mohsen Hariri, **Shouren Wang**, Osama Zafar, Christian Gagné, Vipin Chaudhary[†]
- Links: [\[Paper\]](#)

[10] **CausalGuard: Conformal Inference under Graph Uncertainty**

Under Review

- Authors: Vikash Singh, Weicon Chen, Debargha Ganguly, Yanyan Zhang, Nengbo Wang, Sreehari Sankar, Mohsen Hariri, Alexander Nemecek, Chaoda Song, **Shouren Wang**, Biyou Zhang, Van Yang, Erman Ayday, Jing Ma, Vipin Chaudhary[†]
- Links: [\[Paper\]](#)

[11] **Enhancing Player Enjoyment with a Two-Tier DRL and LLM-Based Agent System for Fighting Games**

Under Review

- Authors: **Shouren Wang**, Zehua Jiang, Fernando Sliva, Sam Earle, Julian Togelius[†]
- Links: [\[Paper\]](#), [\[Repository\]](#)

ACADEMIC SERVICE

Reviewer: Conference on Language Modeling (COLM) 2026

TECHNICAL SKILLS

Languages: Python, Java, MATLAB

Frameworks: PyTorch, Transformers, vLLM, Llama-Factory, Stable-Baselines3, OpenAI Gym, Scikit-Learn

Developer Tools: Git, Slurm, Anaconda, Docker, Singularity, Linux, VS Code

AI Development Tools: Claude Code, OpenClaw, Codex, Cursor

RELEVANT GRADUATE-LEVEL COURSEWORK

- **LLMs & Generative Models:** Large Language Models [CWRU, CSDS 600], Deep Generative Models [CWRU, CSDS 570]
- **Machine Learning:** Advanced Machine Learning [NYU, ECE-GY 7143], Machine Learning [NYU, ECE-GY 6143], Causality and Machine Learning [CWRU, CSDS 452], Data Mining [CWRU, CSDS 435]
- **Core CS & Mathematics:** Data Structures and Algorithms [NYU, ECE-GY 9343], Computing Systems Architecture [NYU, ECE-GY 6913], Probability and Stochastic Processes [NYU, ECE-GY 6303]